

# gemm optimization

JX-Ma

2023/11/24

## 1 实验环境

- 系统:ubuntu20.4
- gcc version:13.1.0
- 优化选项:-O2
- CPU:AMD Ryzen 7 6800H 3.20GHz

## 2 优化函数

### 2.1 无优化

```
void directConvolution_tensor(Tensor<T>& A, Tensor<T>& B, Tensor<T>& C, int64_t s)
{
    for (int64_t i = 0; i < C.num_batch(); ++i){
        for (int64_t j = 0; j < C.num_channel(); ++j){
            for (int64_t m = 0; m < C.num_height(); ++m){
                for (int64_t n = 0; n < C.num_width(); ++n){
                    for (int64_t r = 0; r < B.num_channel(); ++r){
                        for (int64_t u = 0; u < B.num_height(); ++u){
                            for (int64_t v = 0; v < B.num_width(); ++v) {
                                // AI=2/8*3=1/12
                                C(i, j, m, n) += A(i, r, m * s + u, n * s + v) * B(j, r, u, v);
                            }
                        }
                    }
                }
            }
        }
    }
}
```

## 2.2 优化 1

主要优化部分: 利用指针来进行地址偏移寻找数据

```
template<typename T>
void directConvolution_tensor1(Tensor<T>& A, Tensor<T>& B, Tensor<T>& C, int64_t s)
{
    for (int64_t i = 0; i < C.num_batch(); ++i){
        for (int64_t j = 0; j < C.num_channel(); ++j){
            for (int64_t m = 0; m < C.num_height(); ++m){
                for (int64_t n = 0; n < C.num_width(); ++n){
                    for (int64_t r = 0; r < B.num_channel(); ++r){
                        for (int64_t u = 0; u < B.num_height(); ++u){
                            // AI = 6/7*8=3/28
                            AddDot(B.num_width(), &A(i, r, m*s+u, n*s), 1, &B(j, r, u, 0), &C(i, j, m, n));
                            //C.setTensors(i, j, m, n, (A(i, r, m+u, n+v) * B(j, r, u, v)));
                        }
                    }
                }
            }
        }
    }
}

void AddDot(int64_t k, double * A, int intcx, double *B, double *C)
{
    //2*3/8*7
    for(int p = 0; p < k; ++p){
        *C += *A++ * *B++;
    }
}
```

## 2.3 优化 2-分块:1\*3

```
template<typename T>
void directConvolution_tensor2(Tensor<T>& A, Tensor<T>& B, Tensor<T>& C, int64_t s)
{
    for (int64_t i = 0; i < C.num_batch(); ++i){
        for (int64_t j = 0; j < C.num_channel(); ++j){
            for (int64_t m = 0; m < C.num_height(); ++m){
                for (int64_t n = 0; n < C.num_width(); n+=3){
                    for (int64_t r = 0; r < B.num_channel(); r++){
```

```

for (int64_t u = 0; u < B.num_height(); ++u){
    // AI = 6*3/8*(7+6+6)=9/76 //地址偏移次数4*3次
    AddDot(B.num_width(), &A(i, r, m*s+u, n*s), 1, &B(j, r, u, 0), &C(i, j, m, n));
    AddDot(B.num_width(), &A(i, r, m*s+u, n*s+1), 1, &B(j, r, u, 0), &C(i, j, m, n+1));
    AddDot(B.num_width(), &A(i, r, m*s+u, n*s+2), 1, &B(j, r, u, 0), &C(i, j, m, n+2));
}
}
}
}
}
}
}
}

```

## 2.4 优化 2-分块:3\*3

```

template<typename T>
void directConvolution_tensor7(Tensor<T>& A, Tensor<T>& B, Tensor<T>& C, int64_t s) {
for (int64_t i = 0; i < C.num_batch(); ++i) {
for (int64_t j = 0; j < C.num_channel(); ++j) {
for (int64_t m = 0; m < C.num_height(); m+=3) {
for (int64_t n = 0; n < C.num_width(); n += 3) {
for (int64_t r = 0; r < B.num_channel(); r++) {
for (int64_t u = 0; u < B.num_height(); ++u) {
    AddDot(B.num_width(), &A(i, r, m * s + u, n * s), 1,
    &B(j, r, u, 0), &C(i, j, m, n));
    AddDot(B.num_width(), &A(i, r, m * s + u, n * s + 1), 1,
    &B(j, r, u, 0), &C(i, j, m, n + 1));
    AddDot(B.num_width(), &A(i, r, m * s + u, n * s + 2), 1,
    &B(j, r, u, 0), &C(i, j, m, n + 2));

    AddDot(B.num_width(), &A(i, r, m * s + u + 1, n * s), 1,
    &B(j, r, u, 0), &C(i, j, m+1, n));
    AddDot(B.num_width(), &A(i, r, m * s + u + 1, n * s + 1), 1,
    &B(j, r, u, 0), &C(i, j, m+1, n + 1));
    AddDot(B.num_width(), &A(i, r, m * s + u + 1, n * s + 2), 1,
    &B(j, r, u, 0), &C(i, j, m+1, n + 2));

    AddDot(B.num_width(), &A(i, r, m * s + u + 2, n * s), 1,
    &B(j, r, u, 0), &C(i, j, m+2, n));
    AddDot(B.num_width(), &A(i, r, m * s + u + 2, n * s + 1), 1,

```

```

        &B(j, r, u, 0), &C(i, j, m+2, n + 1));
        AddDot(B.num_width(), &A(i, r, m * s + u + 2, n * s + 2), 1,
        &B(j, r, u, 0), &C(i, j, m+2, n + 2));
    }
}
}
}
}
}
}
}

```

## 2.5 优化 3-分块 1\*3

主要优化部分: 充分利用指针偏移,

```

template<typename T>
void directConvolution_tensor3(Tensor<T>& A, Tensor<T>& B, Tensor<T>& C, int64_t s){
    for (int64_t i = 0; i < C.num_batch(); ++i){
        for (int64_t j = 0; j < C.num_channel(); ++j){
            for (int64_t m = 0; m < C.num_height(); ++m){
                for (int64_t n = 0; n < C.num_width(); n+=3){
                    for (int64_t r = 0; r < B.num_channel(); r++){
                        for (int64_t u = 0; u < B.num_height(); ++u){
                            AddDot1x3(B.num_width(), &A(i, r, m*s+u, n*s), 1, &B(j, r, u, 0), &C(i, j, m, n));
                        }
                    }
                }
            }
        }
    }
}

```

```

void AddDot1x3(int64_t k, double * A, int intcx, double *B, double *C)
{
    // A =  $\mathcal{A}(i, r, m*s+u, n*s)$  B =  $\mathcal{B}(j, r, u, 0)$ 
    // AI =  $3*6/19*8=9$ 
    AddDot(k, &(*A++), 1, &(*B), &(*C++));
    AddDot(k, &(*A++), 1, &(*B), &(*C++));
    AddDot(k, &(*A), 1, &(*B), &(*C));
}

```

## 2.6 优化 3-分块 3\*3

```

template<typename T>
void directConvolution_tensor8(Tensor<T>& A, Tensor<T>& B, Tensor<T>& C, int64_t s){
    for (int64_t i = 0; i < C.num_batch(); ++i){
        for (int64_t j = 0; j < C.num_channel(); ++j){
            for (int64_t m = 0; m < C.num_height(); m+=3){
                for (int64_t n = 0; n < C.num_width(); n+=3){
                    for (int64_t r = 0; r < B.num_channel(); r++){
                        for (int64_t u = 0; u < B.num_height(); ++u){
                            AddDot1x3(B.num_width(), &A(i, r, m * s + u, n * s), 1,
                                &B(j, r, u, 0), &C(i, j, m, n));
                            AddDot1x3(B.num_width(), &A(i, r, m * s + u+1, n * s), 1,
                                &B(j, r, u, 0), &C(i, j, m+1, n));
                            AddDot1x3(B.num_width(), &A(i, r, m * s + u+2, n * s), 1,
                                &B(j, r, u, 0), &C(i, j, m+2, n));
                        }
                    }
                }
            }
        }
    }
}

```

## 2.7 优化 4-分块 3\*1

```

// Optimization4 unroll = 3
template<typename T>
void directConvolution_tensor4(Tensor<T>& A, Tensor<T>& B, Tensor<T>& C, int64_t s){
    for (int64_t i = 0; i < C.num_batch(); ++i){
        for (int64_t j = 0; j < C.num_channel(); ++j){
            for (int64_t m = 0; m < C.num_height(); ++m){
                for (int64_t n = 0; n < C.num_width(); n+=3){
                    for (int64_t r = 0; r < B.num_channel(); r++){
                        for (int64_t u = 0; u < B.num_height(); ++u){
                            AddDot1x3_2(B.num_width(), &A(i, r, m*s+u, n*s), 1, &B(j, r, u, 0), &C(i, j, m, n));
                        }
                    }
                }
            }
        }
    }
}

```

```

}
}
}

void AddDot1x3_2(int64_t k, double * A, int intcx, double *B, double *C)
{
    // AddDot(k, &(*A++), 1, &(*B), &(*C++));
    int p = 0;
    double* A1=A;
    double* B1=B;
    for(p = 0; p < k; ++p){
        *C += *A++ * *B++;
    }
    //AddDot(k, &(*A++), 1, &(*B), &(*C++));
    C = C+1; A = A1+1; B = B1;
    for(p = 0; p < k; ++p){
        *C += *A++ * *B++;
    }
    C = C+1; A = A1+1; B = B1;
    //AddDot(k, &(*A), 1, &(*B), &(*C));
    for(p = 0; p < k; ++p){
        *C += *A++ * *B++;
    }
}
}

```

## 2.8 优化 4-分块 3\*3

```

// Optimization4 unroll = 3
template<typename T>
void directConvolution_tensor4(Tensor<T>& A, Tensor<T>& B, Tensor<T>& C, int64_t s){
    for (int64_t i = 0; i < C.num_batch(); ++i){
        for (int64_t j = 0; j < C.num_channel(); ++j){
            for (int64_t m = 0; m < C.num_height(); m+=3){
                for (int64_t n = 0; n < C.num_width(); n+=3){
                    for (int64_t r = 0; r < B.num_channel(); r++){
                        for (int64_t u = 0; u < B.num_height(); ++u){
                            AddDot1x3_2(B.num_width(), &A(i, r, m * s + u, n * s), 1,
                                &B(j, r, u, 0), &C(i, j, m, n));
                            AddDot1x3_2(B.num_width(), &A(i, r, m * s + u + 1, n * s), 1,
                                &B(j, r, u, 0), &C(i, j, m + 1, n));
                            AddDot1x3_2(B.num_width(), &A(i, r, m * s + u + 2, n * s), 1,

```

```

        &B(j, r, u, 0), &C(i, j, m + 2, n));
    }
}
}
}
}
}
}
}

```

## 2.9 优化 5-分块 1\*3

```

// Optimization5 unroll = 3
template<typename T>
void directConvolution_tensor5(Tensor<T>& A, Tensor<T>& B, Tensor<T>& C, int64_t s) {
    for (int64_t i = 0; i < C.num_batch(); ++i) {
        for (int64_t j = 0; j < C.num_channel(); ++j) {
            for (int64_t m = 0; m < C.num_height(); ++m) {
                for (int64_t n = 0; n < C.num_width(); n += 3) {
                    for (int64_t r = 0; r < B.num_channel(); r++) {
                        for (int64_t u = 0; u < B.num_height(); ++u) {
                            // AI = 6*3/8*(7+6+6)=9/76
                            AddDot1x3_3(B.num_width(), &A(i, r, m * s + u, n * s), 1, &B(j, r, u, 0), &C(i, j,
                                m, n));
                        }
                    }
                }
            }
        }
    }
}

void AddDot1x3_3(int64_t k, double* A, int intcx, double* B, double* C)
{
    // AddDot(k, &(*A++), 1, &(*B), &(*C++));
    int p = 0;
    double* C1 = C + 1;
    double* C2 = C1 + 1;
    double* B1 = B;
    double* B2 = B;
    double* A1 = A + 1;
    double* A2 = A1 + 1;
    for (p = 0; p < k; ++p) {

```

```

        *C += *A++ * *B++;
        *C1 += *A1++ * *B1++;
        *C2 += *A2++ * *B2++;
    }
}

```

## 2.10 优化 5-分块 3\*3

```

// Optimization5 unroll = 3
template<typename T>
void directConvolution_tensor10(Tensor<T>& A, Tensor<T>& B, Tensor<T>& C, int64_t s) {
    for (int64_t i = 0; i < C.num_batch(); ++i) {
        for (int64_t j = 0; j < C.num_channel(); ++j) {
            for (int64_t m = 0; m < C.num_height(); m+=3) {
                for (int64_t n = 0; n < C.num_width(); n += 3) {
                    for (int64_t r = 0; r < B.num_channel(); r++) {
                        for (int64_t u = 0; u < B.num_height(); ++u) {
                            AddDot1x3_3(B.num_width(), &A(i, r, m * s + u, n * s), 1,
                                &B(j, r, u, 0), &C(i, j, m, n));
                            AddDot1x3_3(B.num_width(), &A(i, r, m * s + u + 1, n * s), 1,
                                &B(j, r, u, 0), &C(i, j, m + 1, n));
                            AddDot1x3_3(B.num_width(), &A(i, r, m * s + u + 2, n * s), 1,
                                &B(j, r, u, 0), &C(i, j, m + 2, n));
                        }
                    }
                }
            }
        }
    }
}

```

\subsection{优化6-分块1\*3}

\lstset{language=C++}

\begin{lstlisting}

// Optimization6 unroll = 3

template<typename T>

```

void directConvolution_tensor6(Tensor<T>& A, Tensor<T>& B, Tensor<T>& C, int64_t s) {
    for (int64_t i = 0; i < C.num_batch(); ++i) {
        for (int64_t j = 0; j < C.num_channel(); ++j) {

```



```

for (int64_t m = 0; m < C.num_height(); ++m) {
for (int64_t n = 0; n < C.num_width(); n += 3) {
for (int64_t r = 0; r < B.num_channel(); r++) {
for (int64_t u = 0; u < B.num_height(); ++u) {
    AddDot1x3_4(B.num_width(), &A(i, r, m * s + u, n * s), 1, &B(j, r, u, 0), &C(i, j,
    )
  }
}
}
}
}
}
}

void AddDot1x3_4(int64_t k, double* A, int intex, double* B, double* C)
{
    // AddDot(k, &(*A++), 1, &(*B), &(*C++));
    int p = 0;
    double t1, t2, t3;
    t1 = 0.0;
    t2 = 0.0;
    t3 = 0.0;
    double* B1 = B;
    double* B2 = B;
    double* A1 = A + 1;
    double* A2 = A1 + 1;
    for (p = 0; p < k; ++p) {
        t1 += *A++ * *B++;
        t2 += *A1++ * *B1++;
        t3 += *A2++ * *B2++;
    }
    *C++ += t1;
    *C++ += t2;
    *C += t3;
}

```

## 2.11 优化 6-分块 3\*3

```

// Optimization6 unroll = 3
template<typename T>
void directConvolution_tensor11(Tensor<T>& A, Tensor<T>& B, Tensor<T>& C, int64_t s) {
for (int64_t i = 0; i < C.num_batch(); ++i) {

```

```

for (int64_t j = 0; j < C.num_channel(); ++j) {
for (int64_t m = 0; m < C.num_height(); m+=3) {
for (int64_t n = 0; n < C.num_width(); n += 3) {
for (int64_t r = 0; r < B.num_channel(); r++) {
for (int64_t u = 0; u < B.num_height(); ++u) {
    AddDot1x3_4(B.num_width(), &A(i, r, m * s + u, n * s), 1,
        &B(j, r, u, 0), &C(i, j, m, n));
    AddDot1x3_4(B.num_width(), &A(i, r, m * s + u + 1, n * s), 1,
        &B(j, r, u, 0), &C(i, j, m + 1, n));
    AddDot1x3_4(B.num_width(), &A(i, r, m * s + u + 2, n * s), 1,
        &B(j, r, u, 0), &C(i, j, m + 2, n));
    }
}
}
}
}
}
}

```

### 3 实验结果

数据

- input:1x512x8x8
- F: 512x512x3x3
- output:1\*512\*6\*6
- s : 1
- padding: no
- dataType: double

本次实验虽说是 11 种优化算法，实际上只有 6 种，后面五种优化算法分别在分块为 1x3，和分块为 3\*3 做相同的优化

表 1: Gfloat 表

数据	无优化	优化 1	优化 2	优化 3	优化 4	优化 5	优化 6
分块 1x3	0.427695	1.25999	1.49834	2.2886	2.2566	2.31094	2.97743
分块 3x3	0.427695	1.25999	1.55918	2.45546	2.50008	2.55816	3.25192

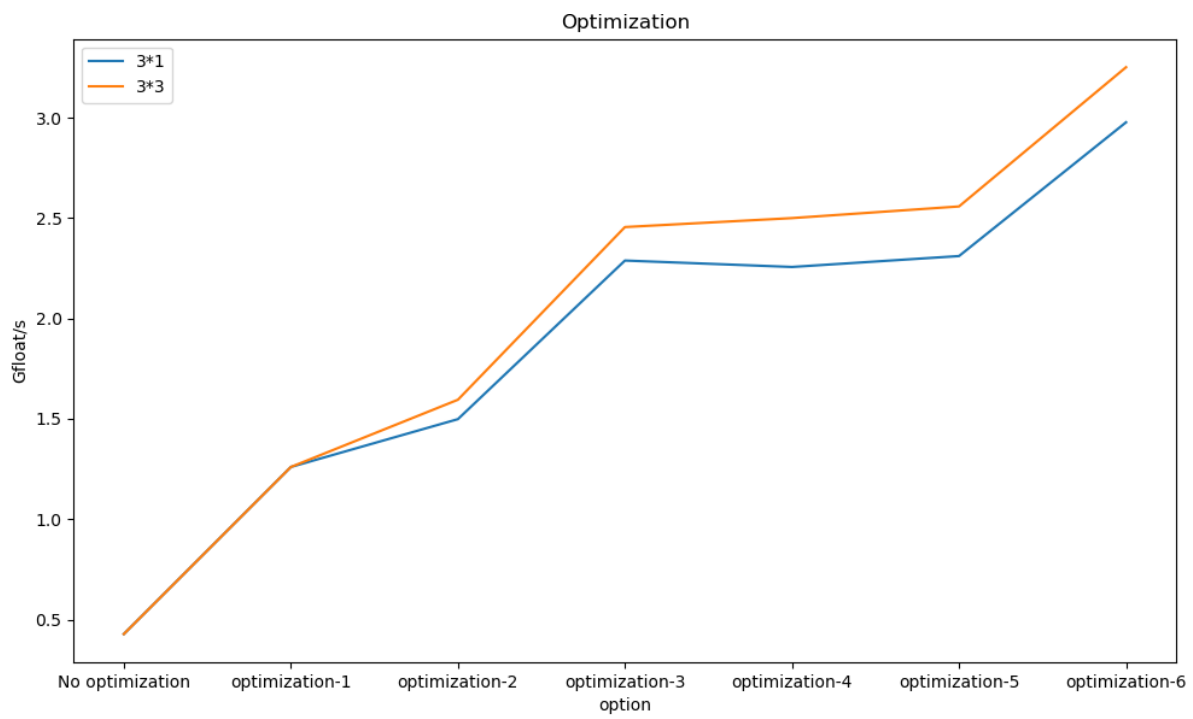


图 1: Gfloat